

SORTING

- Very important problem, most extensively studied
- In many applications, sorting consumes a large proportion of computing time!
- What sorting algorithms have you learned?
 - Bucket Sort and Radix Sort
 - Insertion Sort and Selection Sort
 - Merge Sort
 - Quick Sort
 - Heap Sort
- We will review the above algorithms, do more detailed analysis and prove lower bounds

Bucket Sort

- “Mailroom sort”: allocate a sufficient number of *boxes* – *buckets* – and put each element in the corresponding bucket.
- Works very well only for elements from a small, simple range that is known in advance
 - † e.g. sorting letters by state (by province)
 - † e.g. sorting letters by zip code – we need $26^3 \cdot 10^3$ buckets!!
- Input x_1, x_2, \dots, x_n , $1 \leq x_i \leq m$ and x_i are distinct integers.
Allocate m buckets.
For each i , we put x_i in the bucket corresponding to its value.
Finally, we scan the buckets in order and collect all elements.
- Time and space complexity:
 - time: $O(n + m)$: n for sort n elements and m for final scan
 - space: $O(m)$: 1 unit for each bucket
 - If $m = O(n)$ then this is linear sorting

Radix Sort

- Natural extension of bucket sort.
- We want to reduce the number of buckets (we need more passes).
- Assume that the elements are large integers represented by k digits, and each digit is in the range 0 to $d - 1$.
- We use induction to show the algorithm.

Induction Hypothesis: We know how to sort elements of $< k$ digits

- † Given elements with k digits, we first ignore the most significant digit (left-most digit) and sort the elements according to the rest of the digits by induction!
- † Scan all the elements again and use bucket sort on the most significant digit with d buckets.
- † Collect all the buckets in order.

Example: $n = 10, d = 10, k = 2$

Input: 36, 9, 0, 25, 1, 49, 64, 16, 81, 4

(first pass)		(second pass)	
Bucket	Contents	Bucket	Contents
0	0	0	0, 1, 4, 9
1	1, 81	1	16
2		2	25
3		3	36
4	64, 4	4	49
5	25	5	
6	36, 16	6	64
7		7	
8		8	81
9	9, 49	9	

Append 10 queues of first pass: 0, 1, 81, 64, 4, 25, 36, 16, 9, 49 (input to second pass).

Append 10 queues of second pass: 0, 1, 4, 9, 16, 25, 36, 49, 64, 81

Why does it work?

1. Two elements that are put in different buckets in the LAST step are in the right order
 - do not need induction
 - most significant digits determine the order
2. Two elements having the same most significant digit
 - By induction, they are in right order before the last step.
 - Make sure that elements put in the same bucket REMAIN in the same order
 - using a queue for each bucket
 - appending the d queues at the end of a stage to form one global queue of all elements

(This shows how to use induction to make sure the algorithm is correct.)

Time complexity: $O(kn)$

- Initialize the queues: $O(d)$
- Put n elements into buckets: $O(n)$
- Append d queues: $O(d)$
- Therefore for one pass, total time is $O(n)$

Counting sort can be used to implement Radix Sort.

Counting Sort

Counting sort assumes that each of the n input is an integer in the range of 0 to k , for some k .

The input is in $A[1..n]$ and the output will be in $B[1..n]$.

We also use an array $C[0..k]$ for temporary space.

Counting-Sort(A, B, k)

1. **for** $i = 0$ **to** k
2. **do** $C[i] = 0$
3. **for** $j = 1$ **to** n
4. **do** $C[A[j]] = C[A[j]] + 1$
5. **for** $i = 1$ **to** k
6. **do** $C[i] = C[i] + C[i - 1]$
7. **for** $j = n$ **downto** 1
8. **do** $B[C[A[j]]] = A[j]$
8. $C[A[j]] = C[A[j]] - 1$

- An example that sorting is not based on comparison.
 - Lines 1 to 2 initialize array $C[]$.
 - Lines 3 to 4 calculate the number of integers in the input with value i for each i in the range of 0 and k .
 - Lines 5 to 6 calculate the correct location for the last integer with value i .
 - Lines 7 to 8 sort integers in array $A[]$ and put the result in array $B[]$.
- An important property of counting sort is that it is **stable**.
 - Integers with the same value appear in the output array exactly the same order as they do in the input array.
 - This is important in the application of counting sort.
- When $k = O(n)$, the sort runs in $\Theta(n)$ time.

Insertion Sort

- Assume we can sort $n - 1$ elements, then we can find the right place for the n 'th element and insert it there
- worst case:
 - data movement: $i - 1$ for the i 'th iteration $\implies \Omega(n^2)$
 - comparisons: $\Omega(n \log n)$
- average case:
 - There are i positions where x (i th element) can go.
 - The probability that x belongs to any position is $1/i$.

$$\sum_{j=0}^{i-1} (1/i)j = (1/i)[i(i-1)/2] = (i-1)/2 \quad \textit{ith step}$$

$$A(n) = \sum_{i=2}^n [(i-1)/2] = O(n^2) \quad \textit{total}$$

Selection Sort

- find the maximum element, swap it with the last element.
- data movement: $O(n)$, 1 for each iteration
- comparisons: $O(n^2)$, i for i th largest

Insertion and Selection

Insertion	Data movement	$O(n^2)$
	Comparisons	$O(n \log n)$
Selection	Data movement	$O(n)$
	Comparisons	$O(n^2)$

To improve insertion sort:

Use a data structure that supports search and also insertion, for example AVL trees or red-black trees. These methods require extra space though.

To improve selection sort:

Use a data structure that supports find max and also deletions (e.g. heap sort).

Mergesort

The merge process can be considered as an improvement of insertion sort.

- *Idea:* With the time to insert one element, we can insert many elements.

Let $A = a_1, a_2, \dots, a_n$, $B = b_1, b_2, \dots, b_m$ be sorted.

- We want to insert B into A
- We scan A from the left for the right position for b_1
- We can then continue, without going back, to scan for the right position for b_2 and so on.
- *Data movement:* copy them into a temporary array. Each element moves only once.
- $O(n + m)$ time

Mergesort: Divide-and-conquer sorting

Divide by half $O(1)$

Solve each half recursively $2T(n/2)$

Merge two sorted halves $O(n)$

Time: $T(n) = 2T(n/2) + O(n) \implies O(n \log n)$

QuickSort

Procedure Q-Sort(X , $Left$, $Right$)

begin

if $Left < Right$ then

$Middle = Partition(X, Left, Right);$

$Q-Sort(X, Left, Middle-1);$

$Q-Sort(X, Middle + 1, Right);$

end

Partition

† choose a pivot, x_1

† use two pointers (indices) L and R

Initially, L points to the left end of the array and R points to the right end of the array.

The pointers move in opposite directions.

† *Induction hypothesis: At step k of the partition algorithm, $\text{pivot} \geq x_i$ for $i < L$ and $\text{pivot} < x_j$ for $j > R$.*

1. If $L \leq R$ then

- $x_L \leq \text{pivot}$ then $L \leftarrow L + 1$, or
 $x_R > \text{pivot}$ then $R \leftarrow R - 1$
- $x_L > \text{pivot}$ and $x_R \leq \text{pivot}$ then exchange x_L and x_R , and $L \leftarrow L + 1$, $R \leftarrow R - 1$

2. If $L > R$ exchange x_1 and x_R
(2 = terminate condition)

By induction in step $k + 1$, we can keep induction hypothesis and move either L or R

Consequently, the pointers will eventually meet termination condition

Read textbook pp. 146-148 for another partition algorithm.

How to choose a pivot?

- Choose a random element from the sequence is a good choice
- If the sequence is random, we can just choose the first element
- If we choose another element to be pivot, we can exchange it with the first element, then use our partition algorithm

```

Algorithm Partition(X, left, right);
Input: X (an array)
      left (the left boundary of array X)
      right (the right boundary of array X)
Output: X and middle, such that
       $X[i] \leq X[\text{middle}]$  for all  $i \leq \text{middle}$ ,
       $X[j] > X[\text{middle}]$  for all  $j > \text{middle}$ .
begin
  pivot := X[left]; L := left; R := right;
  while L <= R do
    while X[L] <= pivot and L <= right do
      L := L+1;
    while X[R] > pivot do
      R := R-1;
    if L < R then
      exchange X[L] and X[R]; L := L+1; R := R-1;
  middle := R;
  exchange X[left] and X[middle];
  return middle;
end;

```

Worst case

† The sequence is in the correct order.

$$W(n) = (n - 1) + W(0) + W(n - 1),$$

$$W(0) = W(1) = 1.$$

(($n - 1$) for partition, since the sequence is sorted we will have one empty sequence and one sequence with ($n - 1$) elements.)

Example: 3, 7, 9, 18, 20, 21

$$W(n) = (n - 1) + (n - 2) + 2W(0) + W(n - 2)$$

...

$$= \sum_{i=1}^k (n - i) + kW(0) + W(n - k)$$

$$= \sum_{i=1}^{n-1} (n - i) + (n - 1)W(0) + W(1)$$

$$= \sum_{i=1}^{n-1} i + n = n(n - 1)/2 + n \approx n^2$$

Average Case

- Assume that all keys are distinct. All permutations are equally likely.
- Each x_i has the same probability of being selected as the pivot.
- Running time $T(n)$ of quicksort if the i th smallest element is the pivot is

$$T(n) = (n - 1) + T(i - 1) + T(n - i)$$

$n - 1$ for partition

$T(i - 1)$ for sequence less than pivot

$T(n - i)$ for sequences greater than pivot

- Since x_i has same probability to be pivot the average running time is

$$A(n) = n - 1 + \frac{1}{n} \sum_{i=1}^n [A(i - 1) + A(n - i)],$$

$$A(0) = 0, A(1) = 1.$$

Note

$$\sum_{i=1}^n A(n-i) = A(n-1) + A(n-2) + \cdots + A(0) =$$

$$\sum_{i=1}^n A(i-1).$$

which implies

$$A(n) = (n-1) + \frac{2}{n} \sum_{i=1}^n A(i-1) \quad (1 **)$$

(recurrence with full history)

(1^{**}) involves many $A(i)$'s in $A(n)$. We use a "trick" to reduce this to first order recurrence.

$$nA(n) = n(n-1) + 2 \sum_{i=1}^n A(i-1) \quad (2^{**})$$

$$(n-1)A(n-1) = (n-1)(n-2) + 2 \sum_{i=1}^{n-1} A(i-1) \quad (3^{**})$$

Now subtract (3^{**}) from (2^{**}) :

$$\begin{aligned} nA(n) - (n-1)A(n-1) &= \\ n(n-1) - (n-1)(n-2) + 2A(n-1) &= \\ A(n) = \frac{n+1}{n}A(n-1) + \frac{2(n-1)}{n} &\quad (4^{**}) \end{aligned}$$

Let $B(n) = A(n)/(n+1)$
(Second trick use a substitution.)

$$B(n) = B(n-1) + 2(n-1)/(n+1)n, \quad n > 1.$$

$$B(n) = 2 \sum_{i=2}^n \frac{i-1}{(i+1)i} + B(1)$$

$$B(n) \approx 2 \sum_{i=1}^n \frac{1}{i+1} \approx 2 \sum_{i=1}^n \frac{1}{i}$$

$$\sum_{i=1}^n \frac{1}{i} \approx \ln(n) = \frac{\log n}{\log e}$$

$$B(n) \approx \frac{2}{\log e} \log n \approx 1.4 \log n$$

Conclusion

$$A(n) = 1.4(n+1) \log n$$

$$A(n) = \Theta(n \log n)$$

Space Complexity

- From the appearance of the algorithm, it seems that we do not need any extra space
- However, recursions are implemented by using run-time stacks
Each call: a pair of indices of the array has to be stacked.
- There are at most $n - 1$ calls, there are at most $(n - 1)$ pair of indices to be stacked.
- Space complexity: $O(n)$ (extra space)
- If we use explicit stack, we can guarantee $O(\log n)$ extra space

Improvements of the quicksort algorithm

1. Improve the selection of the pivot
 - choose a random index between L and R .
 - choose the median of x_L , x_R and $x_{(L+R)/2}$

(We need to do extra work, but it is worth it.)
2. Use a simple algorithm for small size.
 - e.g. when size is less than 15, use insertion sort
 - avoid problem of stacking overhead (“choose the base of induction wisely”)
3. Use explicit stacking : avoid overhead of system (run-time) stack
4. Minimize the size of the stack: always stack the larger part first (solve smaller part first).
5. Put pivot into register, for each comparison only one data movement from memory.

Heapsort

- Like selection sort, heapsort is in place
- Like mergesort, heapsort is $O(n \log(n))$
- heapsort combines the better features of the two sorting algorithms.
- heapsort
 - fast sorting algorithm
 - not quite as fast as quicksort
but not much slower
 - unlike quicksort, its performance is guaranteed

Heap Sort

- *Build Heap*
- consider the largest element
 - Swap $A[1], A[n]$
 - $A[n]$ now has correct element
 - rearrange $A[1, \dots, n-1]$ to form a heap (push $A[1]$ down the tree).
- Assume $A[1, \dots, i+1]$ is a heap and $A[i+2], \dots, A[n]$ have correct elements.
 - swap $A[i+1]$ and $A[1]$
 - $A[i+1]$ has now correct element
 - rearrange $A[1, \dots, i]$ to form a heap (push $A[1]$ down the heap).
 - time: $\sum_{i=1}^n \log(i) = O(n \log n)$
(time for transforming a heap to a sorted sequence.)

Time complexity for heap sort

- Heap building
 - $2n$ comparisons
 - n data movements
- Heap sort
 - $2 \sum_{i=2}^n \log i$ comparisons
 - $\sum_{i=2}^n \log i$ data movements
- $\sum_{i=2}^n \log i \leq n \log n - n$
 - $2n + 2 \sum_{i=2}^n \log i \leq 2n \log n$ comparisons.
 - $n + \sum_{i=2}^n \log i \leq n \log n$ data movements.

Lower Bounds for Sorting Problem

- Insertion sort, Selection sort: $O(n^2)$.

Mergesort, heapsort, (quicksort): $O(n \log n)$.

Is it possible to improve it even further?

- Lower bound for a Problem:

A proof that NO Algorithm can solve the problem better.

† Much harder to prove a lower bound for a problem since we have to consider ALL possible algorithms, not just one particular approach.

† We need a model corresponding to an arbitrary (unspecified) algorithm

And a proof that ANY algorithm that fits the model will have a running time higher than the lower bound.

– Example:

We cannot say we will use a special data structure for this problem. Because there may be an algorithm that does not use this data structure and runs faster.

- Decision tree model

- † decision trees model computations that consists of comparisons.

- † Many known lower bound proofs use decision tree model.

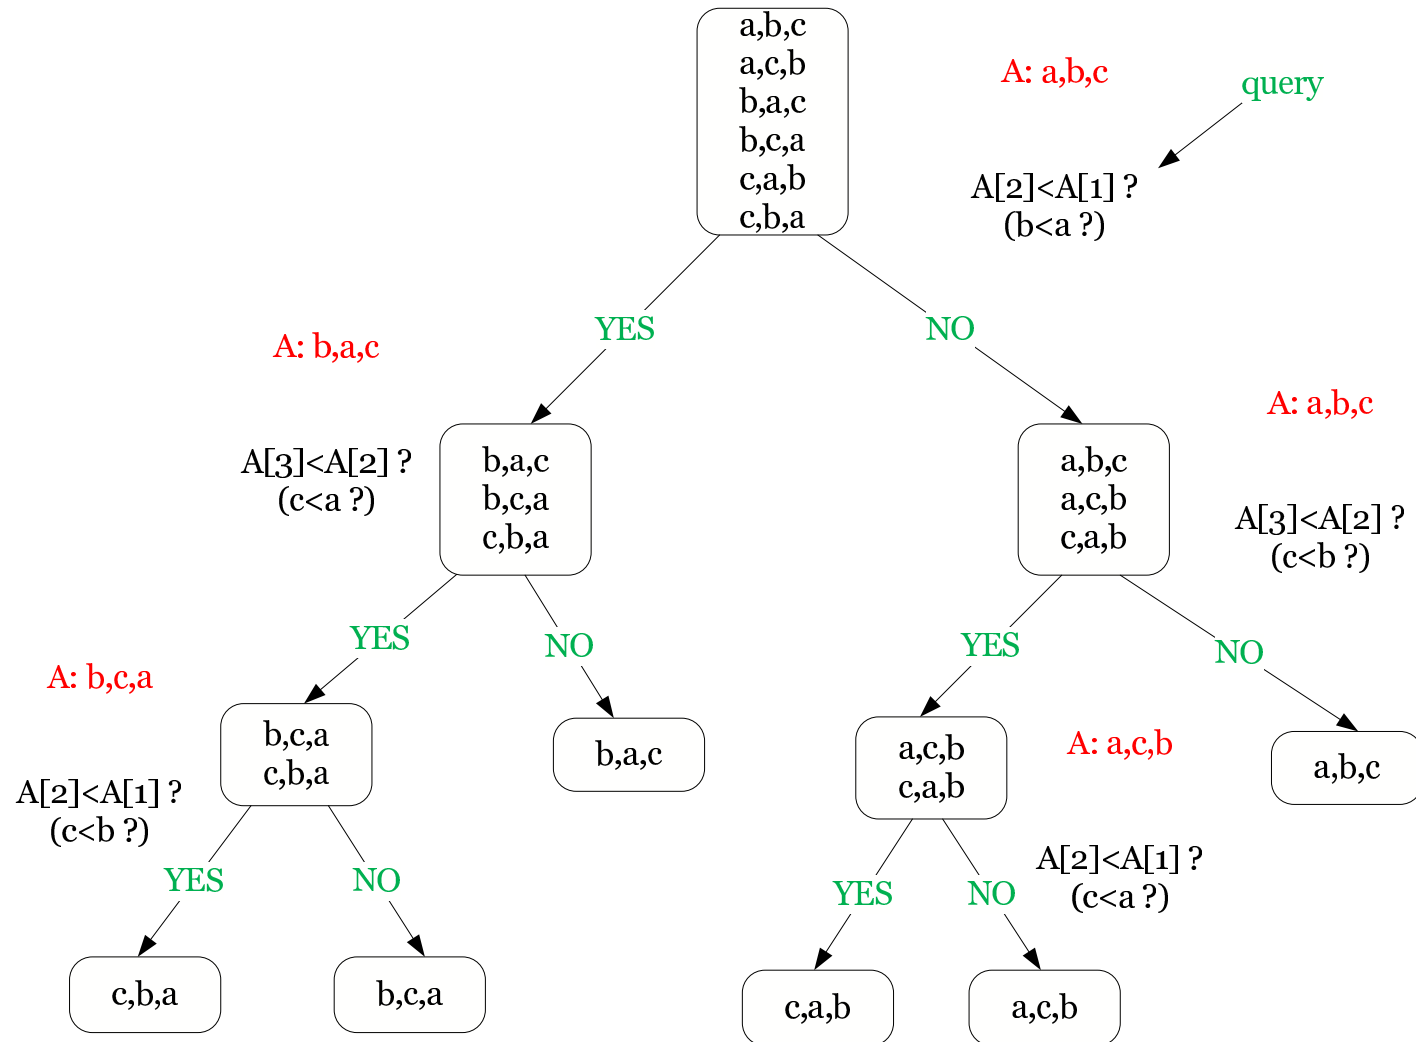
- † As a computation model, decision tree model is weaker than Turing Machine or RAM model.

Decision tree model

- binary trees with two types of nodes:
 internal nodes: two children, *leaves*: no child.
 (Also called two-trees or 2-trees)
- Internal node: associated with a query, the outcome is one of two possibilities. Each one is associated to one of the branches.
- Leaf: associated with a possible output
- Input is a sequence of numbers: x_1, x_2, \dots, x_n
 Computation starts at the root of the tree.
 In each internal node, the query is applied.
 Either go left or go right depending on the result of the query.
- When reaching a leaf, the output associated with the leaf is the output of the computation.
- The worst-case running time of a tree T is the height of T . That is the maximum number of queries required by an input.

The Decision tree for insertion sort with $n = 3$

Input: a, b, c . In array, $A[1, \dots, 3]$



Lower bound for worst case

We want to find the lower bound of the height of a binary tree in terms of number of leaves.

- **Lemma:** Let l be the number of leaves in a binary tree and let h be its height. Then $l \leq 2^h$.

Proof: Induction on h \square .

- Let l and h be as in the Lemma. Then $h \geq \lceil \log l \rceil$. From the Lemma, $l \leq 2^h \implies \log l \leq h \implies h \geq \lceil \log l \rceil$ (since h is an integer).
- A binary tree with $n!$ leaves has a height greater than $n \log n - 1.5n$

$$\begin{aligned} h &\geq \log(n!) \geq \log(n(n-1)(n-2) \dots (\lceil n/2 \rceil)) \\ &\geq \log(\lceil n/2 \rceil^{n/2}) \geq n/2 \log(n/2). \end{aligned}$$

A closer lower bound is

$$h \geq \log(n!) \geq n \log n - 1.5n.$$

Theorem. Every decision tree algorithm for sorting has height $\Omega(n \log n)$.

Proof:

- input for sorting is x_1, x_2, \dots, x_n
- output is a sorted sequence, or is a permutation of input!
(tell us how to rearrange input such that they become sorted.)
- every permutation is a possible output (input can be in any order)
- every permutation of $(1, 2, \dots, n)$ must be represented as an output in the decision tree for sorting.
(Otherwise sorting algorithm is not correct!)
- two different permutations represent two different outputs. They must be associated with different leaves.
- total number of permutations is $n!$
- height of the tree is at least $\log(n!) \geq cn \log n$
- height is $\Omega(n \log n)$. \square

Information-theoretic lower bound

- The lower bound depends only on the amount of information contained in the output.
- It needs to distinguish between $n!$ different outputs; it can only distinguish two possibilities at a time
- Encoding $n!$ possibilities needs $\log(n!)$ bits
- We have not even defined the kind of query we allow
- This lower bound only implies:


NO-COMPARISON-BASED sorting algorithms can be faster than $\Omega(n \log n)$

Lower bound for average behavior

Is it possible to find a comparison algorithm for sorting which has an average behavior better than $n \log n$?

Answer: **NO.**

- epl (external path length) = sum of the length of all paths from the root to a leaf.
- apl (average path length) = epl / (number of leaves)
- *Example:*
epl = $2 + 2 + 1 = 5$ while apl = $5/3 \approx 1.67$



$$\text{epl} = 2 + 2 + 1 = 5$$



$$\text{apl} = 5/3 = 1.67$$

Lemma. Among 2-trees with l leaves, the epl is minimal only if all the leaves are on at most two adjacent levels.

Proof. Suppose we have a 2-tree of height h that has a leaf x at level $k \leq h - 2$.

- choose a node y in level $h - 1$ that is not a leaf
- remove children of y , attach them to x
- the total number of leaves is the same
- net decrease of epl is

$$2h + k - (h - 1 + 2(k + 1)) = h - 1 - k > 0 \quad (\text{since } k \leq h - 2)$$




Lemma. The minimum epl for 2-tree with l leaves is

$$l \lfloor \log l \rfloor + 2(l - 2^{\lfloor \log l \rfloor})$$

Proof.

- From previous Lemma, we can consider only 2-trees of height h and leaves in levels $h - 1$ and h .
- We can transform such a tree into a complete binary tree with possibly some of the right-most leaves removed WITHOUT changing the number of leaves and epl.



- If l is power of two, all leaves are at level $\log l$. $epl = l \log l$
 - If l is not a power of 2, the number leaves at level h is $2(l - 2^{h-1})$
(each node in level $h - 1$ that is not a leaf has two children)
 - $epl = l(h - 1) + 2(l - 2^{h-1}) = l\lfloor \log l \rfloor + 2(l - 2^{\lfloor \log l \rfloor})$
-

Lemma. The average path length in a 2-tree with l leaves is at least $\lfloor \log l \rfloor$.

Proof. The minimum average path length is:

$$\begin{aligned} & \frac{l \lfloor \log l \rfloor + 2(l - 2^{\lfloor \log l \rfloor})}{l} = \\ & = \lfloor \log l \rfloor + 2 \frac{l - 2^{\lfloor \log l \rfloor}}{l} = \\ & = \lfloor \log l \rfloor + \epsilon, \quad 0 \leq \epsilon < 1. \quad \square \end{aligned}$$

$$l/2 < 2^{\lfloor \log l \rfloor} \leq l$$

$$\implies -l/2 > -2^{\lfloor \log l \rfloor} \geq -l$$

$$\implies l - l/2 > l - 2^{\lfloor \log l \rfloor} \geq l - l$$

$$\implies l/2 > l - 2^{\lfloor \log l \rfloor} \geq 0$$

$$\implies 1/2 > [l - 2^{\lfloor \log l \rfloor}]/l \geq 0$$

$$\implies 1 > 2[l - 2^{\lfloor \log l \rfloor}]/l \geq 0$$

Theorem. The average number of comparisons done by an algorithm to sort n numbers by comparison is at least $\lfloor \log n! \rfloor \approx \lfloor n \log n - 1.5n \rfloor$. \square