

## Introduction

- Eight great ideas:
  - Design for Moore's Law.
  - Use abstraction to simplify design.
  - Make common case faster.
  - Parallelism.
  - Pipelining.
  - Prediction.
  - Hierarchy of memories.
  - Dependability via redundancy.

## Chapter 1

- Clock Rate (CR) = 1 / Clock Cycle (CC).
- Effective Clock Cycles per Instruction (CPI) =  $\sum_{I=1}^n = CPI_i \times IC_i$ .
- Instruction Count (IC).
- CPU time = IC x CPI x CC or IC x CPI / CR.

## Chapter 3

- Locality
  - Temporal Locality: if a memory location is referenced it will most likely be referenced again shortly.
  - Spatial Locality: if a memory location is referenced its neighbouring memory location will most likely be referenced shortly.
- Miss types
  - Compulsory/Cold
  - Conflict miss: replacement algorithm keeps mapping the new data to the same block.
    - Ex. 0, 8, 0, 8, 0, 8 if use mod 4 to map blocks then it'll miss every time.
  - Capacity miss: working set is larger than the cache block.
- Average Memory Access Time (AMAT) = Hit time + Miss Rate x Miss Penalty.
- $CPI_{stall} = CPI_{ideal} + \text{Memory-stall cycles}$ .
- Memory-stall cycles = number of access / instruction x miss rate x miss penalty.
- Cache size (C) = #Blocks (B) x #Lines (N) x #Sets (R)
- Block bits (b) =  $\log_2 B$
- Set bits (s) =  $\log_2 R$
- tag bits = number of available bits - b - s
- Types of Cache organization
  - Direct-mapped: One line per set, one block per line. resolves to one block per set.
  - Fully associative: One set with multiple lines, where each line is a block.
  - n-way Set Associative: Multiple sets, with multiple lines, and multiple blocks per line.

- Middle-Order bit indexing is good for spatial locality.
- Write hits
  - Cache-Memory consistent: write to the buffer, only stall when the buffer is full.
  - Cache-Memory inconsistent: write the cache block to the memory when its evicted, write-back. Need a dirty bit to determine if the block needs to be written to memory or not. Can use write buffer to buffer write-backs.

## Chapter 4

- Design Principles:
  - Simplicity favours regularity.
  - Smaller is faster.
  - Make common case faster.
  - Good design demands good compromises.
- MIPS registers
 

◦ \$t0 - \$t9	[8 - 15, 24 - 25]	for temporary values.
◦ \$s0 - \$s7	[16 - 23]	for saved variables.
◦ \$v0 - \$v3	[2 - 3]	for return variables.
◦ \$a0 - \$a3	[4 - 7]	for argument variables.
- Words are 4 bytes each.
- MIPS is Big Endian: most-significant byte at the least address of a word.
- Instruction formats
  - R-format: [ OP | rs | rt | rd | shft | funct ] [6, 5, 5, 5, 5, 6]
  - I-format: [ OP | rs | rt | immediate ] [6, 5, 5, 16]
  - J-format: [ OP | jump target ] [6, 26]
- Operands addressing modes:
  - Immediate.
  - Register addressing.
  - Base addressing: register relative, pseudo-direct.
- Instruction addressing modes:
  - PC-relative: PC + offset, used for branching.
  - Pseudo-direct: jump address, used for jumping.
- Synchronization in MIPS
  - Load Linked: ll rt, offset(rs)
  - Store conditional: sc rt, offset(rs).
  - Succeeds if location not changed since ll. Sets rt to 1.
  - Fails otherwise. Sets rt to 0.

## Chapter 5

- Five stages of the Datapath
  - Instruction Fetch (IF)
  - Instruction Decode (Dec)
  - Arithmetic-Logic Unit (Exec)
  - Memory Access (Mem)
  - Register Write (WB)
- Controllers
  - RegDst = add + sub
  - ALUSrc = ori + lw + sw
  - MemtoReg = lw
  - RegWrite = add + sub + ori + lw
  - MemWrite = sw
  - nPCsel = beq
  - ExtOp = lw + sw
  - ALUctr = add/sub, or

## Chapter 6

- Hazards
  - Structural: attempt to use the same hardware.
  - Data: attempt to use data before its read.
  - Control: attempt to make decision before a condition is evaluated.
- Hazard solutions
  - Forwarding.
  - Delays.
  - Speculation.
- Very Long Instruction Word (VLIW)
- Scheduling/Reordering.
- Unrolling.
- Superscalar
  - Instruction-fetch and issue: fetch, decode and issue instruction to await execution.
  - Instruction-execution: calculate the result, once operands and instruction are ready
  - Instruction-commit: write-back once its safe to do so.
- Instruction and fetch -> functional units -> commit unit.

## Chapter 7

- MESI snooping
  - Modified
  - Exclusive
  - Shared
  - Invalid

- Read Hit
  - No state change
- Read Miss
  - No valid copies
    - Change to E.
  - Another cache has an E
    - Change E to S.
    - Change the rest of the caches to S.
  - Several other caches have S
    - Change local cache to S.
  - One cache has M
    - Share value with caches and Memory.
    - Change all caches to S.
- Write Hit
  - M
    - No change, update local value.
  - E
    - Change E to M.
  - S
    - Change value to M, invalidate other caches.
- Write Miss
  - All caches are invalid
    - Change to M.
  - E , M or S
    - Change to M.
    - Invalidate other caches.
- False Sharing Miss: miss caused, even though the word was not changed.
- True Sharing Miss: miss caused, but the word was changed.
- Types of hardware multithreading
  - Fine-grain: switches threads every instruction.
  - Coarse-grain: switches on costly stalls.
  - Simultaneous Multithreading (SMT)